

# Corpus Explorer

a shiny-driven app for exploratory text analysis



# Goal

- A visual exploratory tool for comparing and contrasting documents in a small-to-medium sized corpus
- High-level document summaries - topic models, significant terms, similar documents
- Easily assess document similarity and clustering visually



# Broad overview

- Preprocessing
- Modeling
- Data structures/conversions ( $\text{Py} \mapsto \text{R}$ )
- Postprocessing
- Final product: exploratory app



# Preprocessing

- Apply custom regex substitution sequence to each document
- Tokenize the results using `Gensim.utils.simple_preprocess`
- Removed tokens with `> max.length` characters (antidisestablishmentarianism: 28)
- Detect bigrams using `Gensim.models.Phrases`
  - based on Mikolov, Tomas et al. "Distributed representations of words and phrases and their compositionality." *Advances in neural information processing systems* 2013.
  - $$\text{score}(w_i, w_j) = \frac{\text{count}(w_i w_j) - \delta}{\text{count}(w_i) \times \text{count}(w_j)}.$$
- This process is iterable: detect n-grams by joining terms to (n-1)-grams, etc.

# Preprocessing (cont.)

- Remove common stop words and terms appearing in  $> p.\text{max}\%$  of documents
- For terms appearing in  $< p.\text{min}\%$  of documents, remove only those occurring  $< \text{min.count}$  times in the corpus
- Save results to log file:


```
Generating corpus-wide dictionary:  
38573 initial terms in the dictionary.
```

```
Computing df's for terms in dictionary.  
70 terms occur in 118 or more documents.  
17934 terms occur in 1 or fewer documents.
```

```
Identifying sparse terms:
```

```
14851 of the rare terms occur 5 or fewer times in any document; removing.
```

```
Removing stop words and sparse terms:  
23569 terms remain in the dictionary after stopwords and rare word removal.
```



# Modeling

- How to summarize documents for easy comparison?
- Significant terms is one way (TF-IDF)
- *Topic Models* are another way: summarize docs as mixtures of “topics” - intuitively, clusters of terms appearing often together
- Why topic models?
  - Low-dimensional representation: 10’s of topics vs 1000’s of terms
  - Semantic coherence: a good quality topic model can produce topics that *make sense* (interpretable)
- Most tried and true flexible topic model: LDA
  - Allows document comparison by fitting one set of topics to the entire corpus

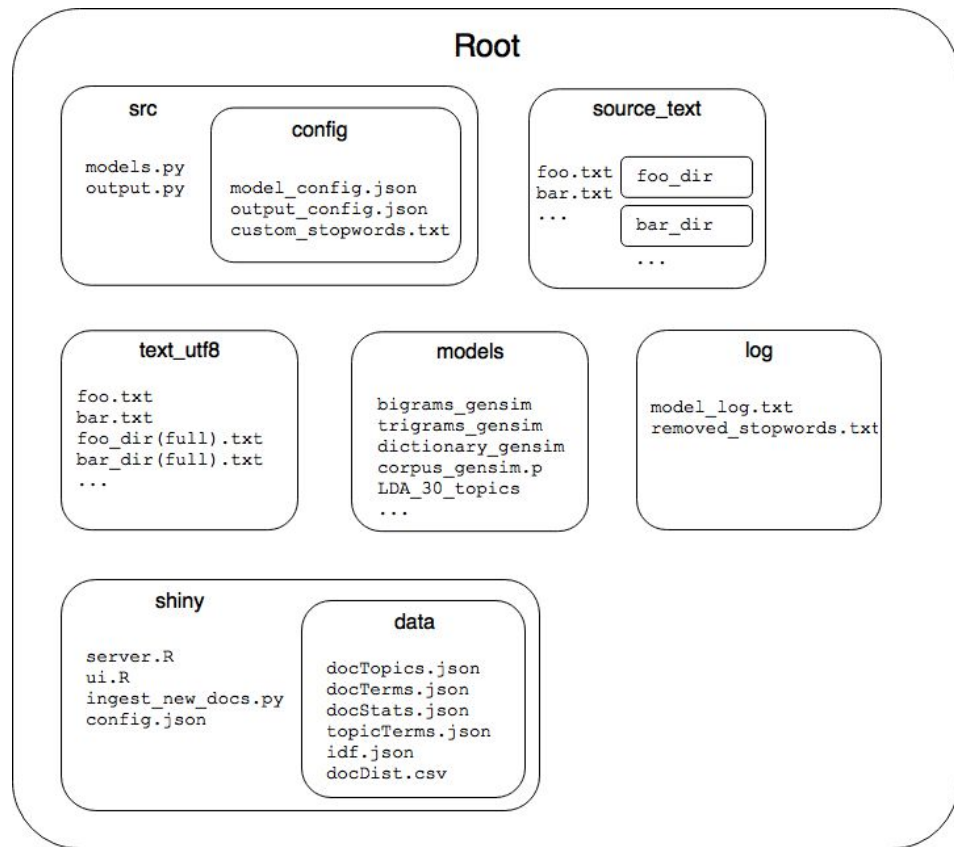


# Modeling (cont.)

- Fit LDA to the corpus using Gensim's implementation
- Since interpretability is important to the end goal, kept topic count small
  - Tried 20, 30, 50, 70, 100, 150 topics
  - Although held-out log-perplexity decreased, topics more redundant, semantically incoherent
  - See: Chang, Jonathan, Sean Gerrish, Chong Wang, Jordan Boyd-Graber, and David M. Blei. "Reading tea leaves: How humans interpret topic models." *Advances in neural information processing systems*. 2009.
- Allowed Gensim to fit an asymmetric topic Dirichlet prior ( $\alpha$ ), generally a better fit than fixed, symmetric prior
  - See: Wallach, Hanna M., David M. Mimno, and Andrew McCallum. "Rethinking LDA: Why priors matter." *Advances in neural information processing systems*. 2009.

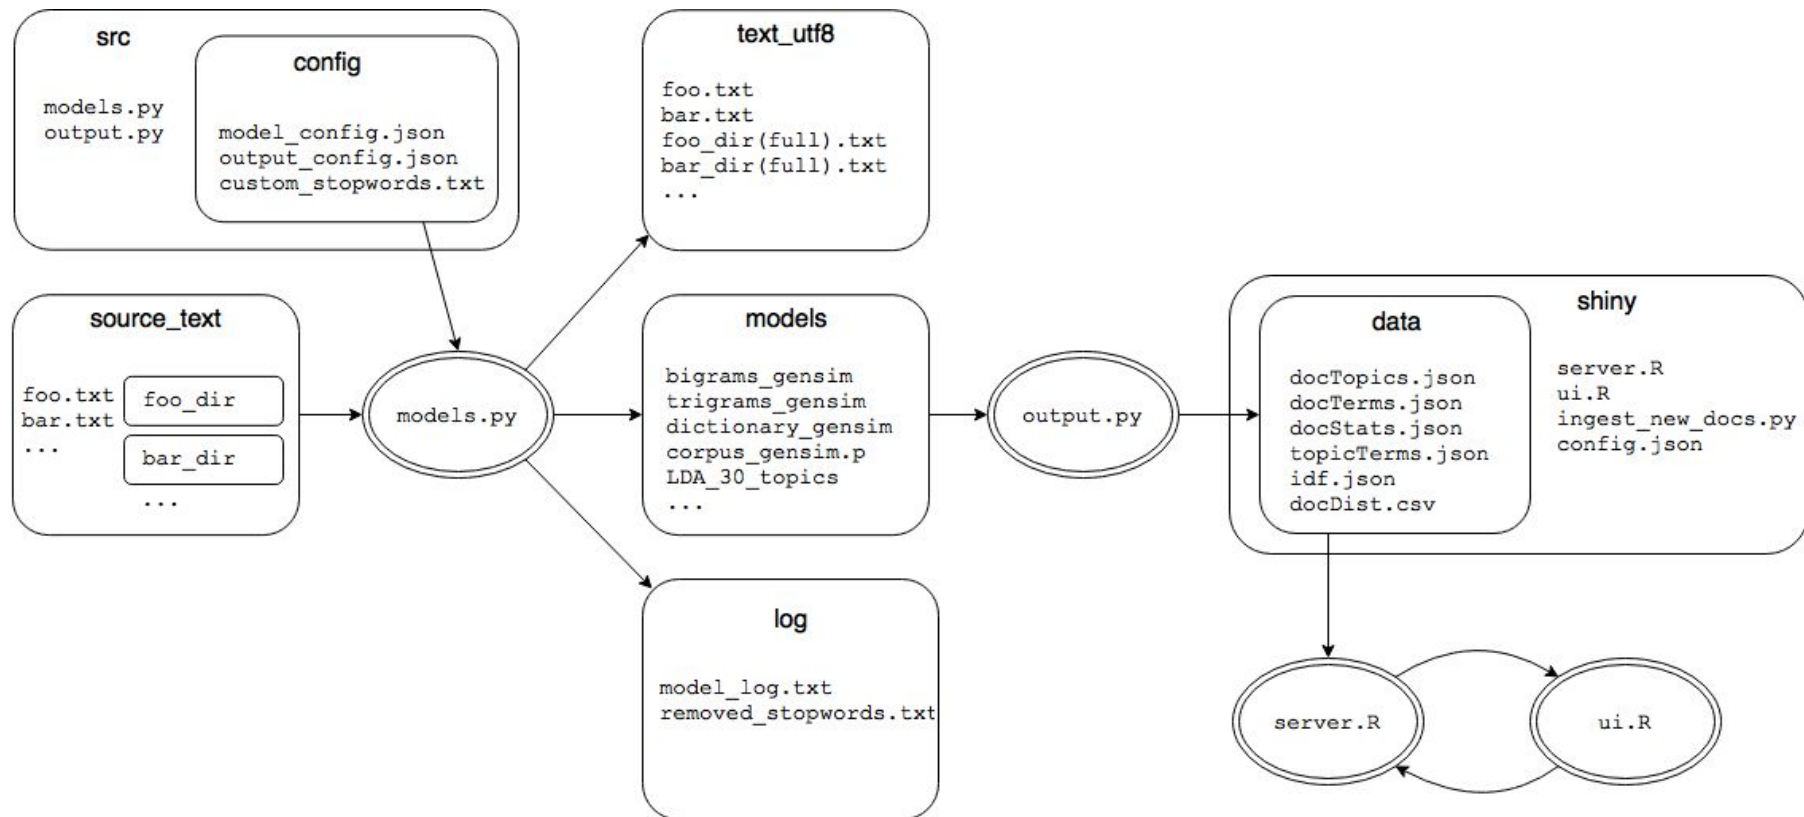


# Directory Structure





# Data Flow




# Data Structures/Formats

- Python needs to talk to R
- Tension between scalability and robustness: smaller  $\Rightarrow$  more fragile
- Ideally for a large corpus, doc-level stats (top terms, topic weights, idf dictionary) would be hashable by document/term name for fast user queries
- Python *dicts* do this
- R *environments* do as well



# Data Structures/Formats

- JSON is an ideal 'glue' format for quickly translating between these types
  - not memory-efficient (redundant), but sparsity helps some:  
27.6MB of text  $\mapsto$  1.6MB of summary .json files (might want to do better for larger corpora)
- With future iterations in mind, JSON is native format for a custom JS app



# Post-processing: document similarity


- Need a good similarity/distance measure
- What's the best notion of distance, given topic distributions? Candidates:
  - Cosine distance (*often outperformed by other metrics; better for tf-idf vectors*)
  - Kullback-Liebler Divergence (*asymmetric and often not defined for sparse distributions*)
  - Hellinger distance - basically a first-order approximation to:
  - Jensen-Shannon Divergence
    - i. rooted in information theory
    - ii. always defined
    - iii.  $\sqrt{\text{J-S}}$  is a proper metric in the mathematical sense - convenient
    - iv. often among the best-performing (see, e.g. *Huang, Anna. "Similarity Measures for Text Document Clustering." 2008.*)

# Post-processing: document clustering

- R reads docDist.csv, a doc-doc distance (J-S) matrix
- Performs agglomerative hierarchical clustering with hclust
- Labels a user-adjustable # of clusters (truncate the hclust merge tree)
- Visualized with a zoomable dendrogram



# Corpus Explorer



# Post-processing: Doc-Topic network

- R reads docTopics.json, a sparse document-topic matrix (specifically, named list of small dataframes, one per doc)
- Builds a bipartite network with documents and topics as nodes
  - edge between doc D and topic T when D is more than p% topic T
  - threshold p is user-adjustable
  - edges are weighted by topic weight; more representative topics end up closer to their docs
  - nodes colored by cluster label from the hierarchical clustering






# Corpus Explorer

Document-Topic Network

Document Dendrogram

Document Stats

Upload New Documents



Documents

ACC-APG RS3.txt

AFSCN CAMMO.txt

AIE-3.txt

AOUSC JMAS OPPS.txt

ASA SOUTHCOM.txt

Topics

0	1	10	11	12	13	14	15	16	17
18	19	2	20	21	22	23	24	25	26
27	28	29	3	4	5	6	7	8	9



# Future Work

- Better handling of OCR mistakes - splitting run-ons
- More comprehensive drill-down UI views - tweaking Shiny at the Javascript level
- Searchability of documents by topic, term, etc.
- Allow user to choose between multiple pre-computed topic models
- Allow user to upload new docs for inference and comparison
- Optimization
  - hierarchical clustering and force networks won't scale to larger corpora - other options?
    - i. spherical or some other variant of k-means
    - ii. local doc-topic network centered on a subset of docs or topics
  - sparse doc-doc distance matrix



Questions?

